

AI Researcher Draft

AI Researcher Agent (fake stages)

Abstract

This draft was generated by the fake writing stage of the AI Researcher Agent. All numbers are bound to aggregate results through `\airval` macros.

1 Introduction

We compare a baseline with a main method and an ablation, following common fine-tuning practice [2, 1].

2 Method

Each experiment from the approved plan was run with three seeds using the task template. Aggregation uses the sample standard deviation and t-based 95% confidence intervals.

3 Results

With 3 seeds each, main reaches a mean score of 78.0% compared with 69.4% for baseline.

The mean difference is 8.6 percentage points (main higher than baseline; Welch $p = 0.0017$, reported for reference only).

With 3 seeds each, main reaches a mean score of 78.0% compared with 73.7% for ablation.

The mean difference is 4.3 percentage points (main higher than ablation; Welch $p = 0.0343$, reported for reference only).

4 Conclusion

These results come from a small number of seeds on a single task and should be read as preliminary.

References

- [1] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *NAACL*, 2019.
- [2] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. In *ICLR*, 2022.